

# DS630(MACHINE LEARNING) FINAL PROJECT

## -SIMI SUDHAKARAN

In [21]:

```
# IMPORT THE NECESSARY LIBRARIES

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

from sklearn.preprocessing import StandardScaler
from sklearn.decomposition import PCA
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score, confusion_matrix, precision_recall_f
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.ensemble import AdaBoostClassifier
```

In [4]:

```
# IMPORT THE DATASET

df = pd.read_excel('CovidDeathUS2021.xlsx')
df.head()
```

Out[4]:

	PEOPLE_POSITIVE_CASES_COUNT	COUNTY_NAME	PROVINCE_STATE_NAME	REPORT_DATE	C
0	9271	Marshall	Alabama	2021-01-01	
1	1517	Barbour	Alabama	2021-01-01	
2	1522	Butler	Alabama	2021-01-01	
3	2418	Clarke	Alabama	2021-01-01	
4	1756	Pickens	Alabama	2021-01-01	

In [5]:

```
# WHOLE INFORMATION OF THE DATASET
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 248083 entries, 0 to 248082
Data columns (total 15 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   PEOPLE_POSITIVE_CASES_COUNT          248083 non-null  int64
1   COUNTY_NAME                          248083 non-null  object
2   PROVINCE_STATE_NAME                  248083 non-null  object
```

```

3  REPORT_DATE                248083 non-null  datetime64[ns]
4  CONTINENT_NAME             248083 non-null  object
5  DATA_SOURCE_NAME          248083 non-null  object
6  PEOPLE_DEATH_NEW_COUNT     248083 non-null  int64
7  COUNTY_FIPS_NUMBER         244654 non-null  float64
8  COUNTRY_ALPHA_3_CODE       248083 non-null  object
9  COUNTRY_SHORT_NAME         248083 non-null  object
10 COUNTRY_ALPHA_2_CODE       248083 non-null  object
11 PEOPLE_POSITIVE_NEW_CASES_COUNT 248083 non-null  int64
12 PEOPLE_DEATH_COUNT         248083 non-null  int64
13 IS_COVID_PRESENT           248083 non-null  bool
14 ARE_THERE_NEW_COVID_DEATHS  248083 non-null  bool
dtypes: bool(2), datetime64[ns](1), float64(1), int64(4), object(7)
memory usage: 25.1+ MB

```

```

In [6]: # DROP IRRELEVANT COLUMNS
df_filtered = df.drop(columns=['COUNTY_NAME', 'PROVINCE_STATE_NAME', 'REPORT_DATE',
                                'COUNTY_FIPS_NUMBER', 'COUNTRY_ALPHA_3_CODE', 'COUNTRY_SHORT_NAME', 'COUNTY_ALPHA_2_CODE'])

```

```

In [7]: # 5 POINT SUMMARY OF THE DATASET
df_filtered.describe()

```

```

Out[7]:

```

	PEOPLE_POSITIVE_CASES_COUNT	PEOPLE_DEATH_NEW_COUNT	PEOPLE_POSITIVE_NEW_C
count	2.480830e+05	248083.000000	2
mean	8.834160e+03	0.562425	
std	3.862826e+04	5.082005	
min	0.000000e+00	-1009.000000	
25%	8.520000e+02	0.000000	
50%	2.062000e+03	0.000000	
75%	5.470000e+03	0.000000	
max	1.235783e+06	930.000000	

```

In [8]: # NEW AVERAGE DEATH COUNT PER DAY
df_filtered['PEOPLE_DEATH_NEW_COUNT'].mean()

```

```

Out[8]: 0.5624246723878702

```

```

In [9]: # TOTAL AVERAGE DEATH COUNT
df_filtered['PEOPLE_DEATH_COUNT'].mean()

```

```

Out[9]: 162.0288290612416

```

```

In [10]: # CORRELATION BETWEEN THE DIFFERENT FEATURES
df_filtered.corr()

```

```

Out[10]:

```

	PEOPLE_POSITIVE_CASES_COUNT	PEOPLE_DEATH_NEW_COUNT
--	-----------------------------	------------------------

	PEOPLE_POSITIVE_CASES_COUNT	PEOPLE_DEATH_NEW_C
PEOPLE_POSITIVE_CASES_COUNT	1.000000	0.5
PEOPLE_DEATH_NEW_COUNT	0.544253	1.0
PEOPLE_POSITIVE_NEW_CASES_COUNT	0.607500	0.6
PEOPLE_DEATH_COUNT	0.908214	0.4
IS_COVID_PRESENT	0.097775	0.0
ARE_THERE_NEW_COVID_DEATHS	0.237220	0.2

```
In [11]: # LOTTING THE FEATURE CORRELATION
correlation = df_filtered.corr()
plt.figure(figsize=(10,10))
sns.heatmap(correlation, vmax=1, square = True, annot=True, cmap = 'viridis')
plt.title("US Covid19(2021) : Correlation between the different features")
```

```
Out[11]: Text(0.5, 1.0, 'US Covid19(2021) : Correlation between the different features')
```



```
In [12]: # INFORMATION FOR THE FILTERED DATASET
df_filtered.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 248083 entries, 0 to 248082
Data columns (total 6 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   PEOPLE_POSITIVE_CASES_COUNT           248083 non-null int64
1   PEOPLE_DEATH_NEW_COUNT                 248083 non-null int64
2   PEOPLE_POSITIVE_NEW_CASES_COUNT        248083 non-null int64
3   PEOPLE_DEATH_COUNT                     248083 non-null int64
4   IS_COVID_PRESENT                       248083 non-null bool
5   ARE_THERE_NEW_COVID_DEATHS             248083 non-null bool
dtypes: bool(2), int64(4)
memory usage: 8.0 MB
```

```
In [13]: # CONVERT THE CATEGORICAL VARIABLES TO DUMMY VARIABLES
df_filtered['ARE_THERE_NEW_COVID_DEATHS'] = df_filtered['ARE_THERE_NEW_COVID_DEA
```

```
df_filtered['IS_COVID_PRESENT'] = df_filtered['IS_COVID_PRESENT'].astype(int)
print(df_filtered)
```

	PEOPLE_POSITIVE_CASES_COUNT	PEOPLE_DEATH_NEW_COUNT	\
0	9271	0	
1	1517	1	
2	1522	0	
3	2418	1	
4	1756	0	
...	...	...	
248078	369	0	
248079	1405	0	
248080	4329	0	
248081	5179	0	
248082	3157	0	

  

	PEOPLE_POSITIVE_NEW_CASES_COUNT	PEOPLE_DEATH_COUNT	IS_COVID_PRESENT	\
0	82	86	1	
1	3	33	1	
2	14	45	1	
3	28	26	1	
4	22	26	1	
...	...	...	...	
248078	0	3	0	
248079	2	12	1	
248080	27	12	1	
248081	0	84	0	
248082	5	31	1	

  

	ARE_THERE_NEW_COVID_DEATHS
0	0
1	1
2	0
3	1
4	0
...	...
248078	0
248079	0
248080	0
248081	0
248082	0

[248083 rows x 6 columns]

In [14]:

```
# ASSIGN THE TARGET VARIABLE(y) AND FEATURE VARIABLE(x)
x = df_filtered.iloc[:,0:3]
y = df_filtered.iloc[:, 5]
```

In [15]:

```
print(x)
```

	PEOPLE_POSITIVE_CASES_COUNT	PEOPLE_DEATH_NEW_COUNT	\
0	9271	0	
1	1517	1	
2	1522	0	
3	2418	1	
4	1756	0	
...	...	...	
248078	369	0	
248079	1405	0	
248080	4329	0	
248081	5179	0	

248082

3157

0

```

    PEOPLE_POSITIVE_NEW_CASES_COUNT
0                                82
1                                 3
2                                14
3                                28
4                                22
...                               ...
248078                           0
248079                           2
248080                           27
248081                           0
248082                           5

```

[248083 rows x 3 columns]

In [16]:

```
print(y)
```

```

0      0
1      1
2      0
3      1
4      0
..
248078  0
248079  0
248080  0
248081  0
248082  0
Name: ARE_THERE_NEW_COVID_DEATHS, Length: 248083, dtype: int64

```

In [12]:

```

#SCALE THE DATA BETWEEN 1 AND 0

standardScaler = StandardScaler()
x = standardScaler.fit_transform(x)
x

```

```

Out[12]: array([[ 0.01130884, -0.11067006,  0.24204903],
                [-0.18942544,  0.08610307, -0.13394795],
                [-0.189296   , -0.11067006, -0.08159394],
                ...,
                [-0.11662884, -0.11067006, -0.01972102],
                [-0.09462418, -0.11067006, -0.14822631],
                [-0.14696939, -0.11067006, -0.12442904]])

```

In [13]:

```

# PERFORM PCA ON THE DATASET
pca = PCA()
pca.fit_transform(x)

```

```

Out[13]: array([[ 0.08639834,  0.06838149,  0.24253409],
                [-0.13758679, -0.18979688, -0.07927678],
                [-0.21908273, -0.06061766,  0.05537745],
                ...,
                [-0.14119856, -0.01061588,  0.07867233],
                [-0.204776   ,  0.01323029, -0.03262496],
                [-0.22040832, -0.02682995,  0.00554753]])

```

In [17]:

```

# SPLITTING THE DATASET TO TEST AND TRAIN
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.2, random_state=42)

```

```
In [18]: print('x_train: ', x_train.shape)
print('x_test: ', x_test.shape)
print('y_train: ', y_train.shape)
print('y_test: ', y_test.shape)
```

```
x_train: (198466, 3)
x_test: (49617, 3)
y_train: (198466,)
y_test: (49617,)
```

```
In [19]: # FUNCTION TO PRINT ALL THE MODEL ACCURACIES THAT WAS CALCULATED DURING GRID SEARCH
def print_results(results):
    print('BEST PARAMS: {}'.format(results.best_params_))

    means = results.cv_results_['mean_test_score']
    stds = results.cv_results_['std_test_score']
    for mean, std, params in zip(means, stds, results.cv_results_['params']):
        print('{} (+/-{}) for {}'.format(round(mean, 3), round(std * 2, 3), para
```

## LOGISTIC REGRESSION CLASSIFIER

```
In [60]: # RUNNING THE LOGISTIC REGRESSION CLASSIFIER USING GRID SEARCH
lr_classifier = LogisticRegression()
parameters = {
    'C' : [0.001] #Inverse regularization parameter
}
cv = GridSearchCV(lr_classifier, parameters, cv = 5) #cv is for k fold cross val
cv.fit(x_train, y_train)
```

```
Out[60]: GridSearchCV(cv=5, estimator=LogisticRegression(), param_grid={'C': [0.001]})
```

```
In [61]: # GIVES US THE HYPERPARAMETER SETTINGS WITH BEST ACCURACY SCORE
cv.best_estimator_
```

```
Out[61]: LogisticRegression(C=0.001)
```

```
In [62]: # PREDICT THE TARGET VALUE USING THE FITTED LOGISTIC REGRESSION CLASSIFIER
y_pred = cv.predict(x_test)
print('Logistic Regression: Accuracy Score - ', accuracy_score(y_test, y_pred))
```

```
Logistic Regression: Accuracy Score - 0.8894733659834331
```

## K-NEAREST NEIGHBOR CLASSIFIER

```
In [49]: # RUN THE K NEAREST NEIGHBOR CLASSIFIER ON THE DATASET
knn_classifier = KNeighborsClassifier()
parameters = {
    'n_neighbors':[5] #no of neighbors
}
```

```
cv_knn = GridSearchCV(knn_classifier, parameters, cv = 5, scoring='accuracy', n_
cv_knn.fit(x_train, y_train)
```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 2 out of 5 | elapsed: 58.1s remaining: 1.5min
[Parallel(n_jobs=-1)]: Done 5 out of 5 | elapsed: 59.7s finished
```

```
Out[49]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(), n_jobs=-1,
param_grid={'n_neighbors': [5]}, scoring='accuracy', verbose=1)
```

```
In [50]: #PREDICT THE ACCURACY SCORE USING THE FITTED KNN CLASSIFIER
y_pred = cv_knn.predict(x_test)
print('K-Nearest Neighbor: Accuracy Score - ', accuracy_score(y_test, y_pred))
```

K-Nearest Neighbor: Accuracy Score - 0.9995364492008787

## KNN on the non-PCA dataset

```
In [22]: # RUN THE KNN CLASSIFIER ON NON SCALED DATASET TO CHECK FOR BETTER ACCURACY
knn_classifier = KNeighborsClassifier()
parameters = {
    'n_neighbors':[5]
}
cv_knn = GridSearchCV(knn_classifier, parameters, cv = 5, scoring='accuracy', n_
cv_knn.fit(x_train, y_train)
```

Fitting 5 folds for each of 1 candidates, totalling 5 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 2 out of 5 | elapsed: 5.7s remaining: 8.6s
[Parallel(n_jobs=-1)]: Done 5 out of 5 | elapsed: 5.9s finished
```

```
Out[22]: GridSearchCV(cv=5, estimator=KNeighborsClassifier(), n_jobs=-1,
param_grid={'n_neighbors': [5]}, scoring='accuracy', verbose=1)
```

```
In [28]: print_results(cv_knn)
```

BEST PARAMS: {'n\_neighbors': 5}

0.852 (+/-0.003) for {'n\_neighbors': 5}

```
In [32]: # GENERATE A CLASSIFICATION REPORT TO CHECK FOR PRECISION, ACCURACY AND RECALL V
y_pred = cv_knn.predict(x_test)
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.88	0.95	0.92	41118
1	0.62	0.38	0.47	8499
accuracy			0.85	49617
macro avg	0.75	0.66	0.69	49617
weighted avg	0.84	0.85	0.84	49617

```
In [ ]:
```



## KERNEL SUPPORT VECTOR MACHINE CLASSIFIER

```
In [84]: # RUN THE DATASET ON KERNEL(RBF) SUPPORT VECTOR MACHINE
svc_classifier = SVC(kernel = 'rbf', random_state = 0)
parameter = {
    'C' : [0.001, 0.01, 0.1, 1, 10, 100]
}
cv_svc = GridSearchCV(svc_classifier, parameter , cv=5)
cv_svc.fit(x_train,y_train)
```

```
Out[84]: GridSearchCV(cv=5, estimator=SVC(random_state=0),
    param_grid={'C': [0.001, 0.01, 0.1, 1, 10, 100]})
```

```
In [87]: y_pred = cv_svc.predict(x_test)
print('SUPPORT VECTOR MACHINE CLASSIFIER: Accuracy Score - ', accuracy_score(y_t

SUPPORT VECTOR MACHINE CLASSIFIER: Accuracy Score - 0.9998186105568656
```

### SVM(kernel) on the non-PCA dataset

```
In [ ]: from sklearn.svm import SVC

svc_classifier = SVC(random_state = 0)
parameter = {
    'kernel' : ['linear', 'rbf'],
    'C' : [0.001, 0.01, 0.1, 1, 10, 100]
}
cv_svc = GridSearchCV(svc_classifier, parameter , cv=5, scoring='accuracy', verb
cv_svc.fit(x_train,y_train)
```

Fitting 5 folds for each of 12 candidates, totalling 60 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

```
In [ ]: print_results(cv_svc)
```

```
In [ ]: y_pred = cv_svc.predict(x_test)
print(classification_report(y_test,y_pred))
```

```
In [ ]:
```

## DECISION TREE CLASSIFIER

```
In [36]: # RUN THE DECISION TREE CLASSIFIER ON THE TEST AND TRAIN DATASET
parameters = {
    'criterion':['entropy'],
    'splitter':['random'],
    'max_depth':[15],
    'min_samples_split':[4],
    'min_samples_leaf':list(range(1, 5)),
}
```

```
dt_classifier = DecisionTreeClassifier(random_state = 42)
cv_dt = GridSearchCV(dt_classifier, parameters, scoring='accuracy', n_jobs=-1, ve
cv_dt.fit(x_train, y_train)
```

Fitting 3 folds for each of 4 candidates, totalling 12 fits

```
[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 10 out of 12 | elapsed: 0.4s remaining: 0.1s
[Parallel(n_jobs=-1)]: Done 12 out of 12 | elapsed: 0.5s finished
```

```
Out[36]: GridSearchCV(cv=3, estimator=DecisionTreeClassifier(random_state=42), n_jobs=-1,
                    param_grid={'criterion': ['entropy'], 'max_depth': [15],
                                'min_samples_leaf': [1, 2, 3, 4],
                                'min_samples_split': [4], 'splitter': ['random']},
                    scoring='accuracy', verbose=1)
```

```
In [38]: cv_dt.best_estimator_
```

```
Out[38]: DecisionTreeClassifier(criterion='entropy', max_depth=15, min_samples_leaf=2,
                                min_samples_split=4, random_state=42, splitter='random')
```

```
In [39]: y_pred = cv_dt.predict(x_test)
         print('DECISION TREE CLASSIFIER: Accuracy Score - ', accuracy_score(y_test, y_pr
```

```
DECISION TREE CLASSIFIER: Accuracy Score - 0.941874760666707
```

## Decion Tree Classifier on the non-PCA dataset

```
In [ ]: # RUN THE DECISION TREE CLASSIFIER ON THE TEST AND TRAIN DATASET
        parameters = {
            'criterion': ["entropy"],
            'splitter': ["random"],
            'max_depth': [15],
            'min_samples_split': [4],
            'min_samples_leaf': list(range(1, 5)),
        }

        dt_classifier = DecisionTreeClassifier(random_state = 42)
        cv_dt = GridSearchCV(dt_classifier, parameters, scoring='accuracy', n_jobs=-1, ve
        cv_dt.fit(x_train, y_train)
```

```
In [ ]: print_results(cv_dt)
```

```
In [ ]: y_pred = cv_dt.predict(x_test)
        print(classification_report(y_test, y_pred))
```

## RANDOM FOREST CLASSIFIER

```
In [41]: # RUN THE RANDOM FOREST CLASSIFIER ON THE TEST AND TRAIN DATASET
        parameters = {
            'n_estimators' : [100, 500, 1000, 2000],
            'max_features' : ['auto', 'sqrt'],
            'max_depth' : [2, 3, 5],
            'min_samples_split' : [2, 5, 10],
```

```

    'min_samples_leaf' : [1,2,4,10],
}

rf_classifier = RandomForestClassifier(random_state=42)
cv_rf = GridSearchCV(rf_classifier, parameters, cv=5, verbose =1, n_jobs=-1)
cv_rf.fit(x_train, y_train)

```

Fitting 5 folds for each of 288 candidates, totalling 1440 fits

```

[Parallel(n_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.
[Parallel(n_jobs=-1)]: Done 34 tasks      | elapsed: 3.1min
[Parallel(n_jobs=-1)]: Done 184 tasks    | elapsed: 18.2min
[Parallel(n_jobs=-1)]: Done 434 tasks    | elapsed: 45.7min
[Parallel(n_jobs=-1)]: Done 784 tasks    | elapsed: 90.0min
[Parallel(n_jobs=-1)]: Done 1234 tasks   | elapsed: 148.8min
[Parallel(n_jobs=-1)]: Done 1440 out of 1440 | elapsed: 175.9min finished

```

```

Out[41]: GridSearchCV(cv=5, estimator=RandomForestClassifier(random_state=42), n_jobs=-1,
                    param_grid={'max_depth': [2, 3, 5],
                                'max_features': ['auto', 'sqrt'],
                                'min_samples_leaf': [1, 2, 4, 10],
                                'min_samples_split': [2, 5, 10],
                                'n_estimators': [100, 500, 1000, 2000]},
                    verbose=1)

```

```

In [43]: y_pred = cv_rf.predict(x_test)
         print('DECISION TREE CLASSIFIER: Accuracy Score - ', accuracy_score(y_test, y_pr

DECISION TREE CLASSIFIER: Accuracy Score - 1.0

```

```

In [ ]: # Random Forest Classifier on a non-PCA dataset

```

```

In [ ]: # RUN THE RANDOM FOREST CLASSIFIER ON THE NON-SCALED DATASET
        parameters = {
            'n_estimators' : [100,500,1000,2000],
            'max_features' : ['auto', 'sqrt'],
            'max_depth' : [2,3,5],
            'min_samples_split' : [2,5,10],
            'min_samples_leaf' : [1,2,4,10],
        }

        rf_classifier = RandomForestClassifier(random_state=42)
        cv_rf = GridSearchCV(rf_classifier, parameters, cv=5, verbose =1, n_jobs=-1)
        cv_rf.fit(x_train, y_train)

```

```

In [ ]: print_result(cv_rf)

```

```

In [ ]: y_pred = cv_rf.predict(x_test)
         print(classification_report(y_test,y_pred))

```

## ENSEMBLING METHOD: ADABOOST ALGORITHM

```

In [77]: # USE THE ENSEMBLING TECHNIQUES TO BOOST THE ACCURACY SCORE
         adaboost_clf = AdaBoostClassifier(random_state=42)

```

```

parameters = {
    'learning_rate' : [0.1,1,0.01,0.5],
    'n_estimators' : [100,120,140,160,180, 200]
}
cv_adaboost = GridSearchCV(adaboost_clf, parameters, cv = 5, n_jobs=-1, verbose=
cv_adaboost.fit(x_train, y_train)

```

Fitting 5 folds for each of 24 candidates, totalling 120 fits

[Parallel(n\_jobs=-1)]: Using backend LokyBackend with 8 concurrent workers.

[Parallel(n\_jobs=-1)]: Done 34 tasks | elapsed: 3.0s

[Parallel(n\_jobs=-1)]: Done 120 out of 120 | elapsed: 5.7s finished

```

Out[77]: GridSearchCV(cv=5, estimator=AdaBoostClassifier(random_state=42), n_jobs=-1,
    param_grid={'learning_rate': [0.1, 1, 0.01, 0.5],
    'n_estimators': [100, 120, 140, 160, 180, 200]},
    verbose=1)

```

```

In [78]: y_pred = cv_adaboost.predict(x_test)
    print('ADABOOST CLASSIFIER: Accuracy Score - ', accuracy_score(y_test, y_pred))

```

ADABOOST CLASSIFIER: Accuracy Score - 1.0

```

In [ ]: # Adaboost on a non-PCA dataset

```

```

In [ ]: # USING THE ADABOOST ENSEMBLE TECHNIQUE TO BOOST THE ACCURACY SCORE USING THE NO
    adaboost_clf = AdaBoostClassifier(random_state=42)
    parameters = {
        'learning_rate' : [0.1,1,0.01,0.5],
        'n_estimators' : [100,120,140,160,180, 200]
    }
    cv_adaboost = GridSearchCV(adaboost_clf, parameters, cv = 5, n_jobs=-1, verbose=
    cv_adaboost.fit(x_train, y_train)

```

```

In [ ]: print_result(cv_adaboost)

```

```

In [ ]: y_pred = cv_adaboost.predict(x_test)
    print(classification_report(y_test,y_pred))

```

```

In [ ]:

```

```

In [ ]:

```